

Deep Learning Latency Attacks and Defenses: A Cross-Domain Survey

ZONGHUA GU, Hofstra University, USA

ZEYU GAO, Umeå University, Sweden

AMIN SAREMI, Umeå University, Sweden

Adversarial machine learning has historically focused on *integrity*: protecting the model against attacks that cause it to produce wrong outputs. A complementary and increasingly consequential axis is *availability*, where the attacker’s objective is not to corrupt a prediction but to make the system fail to respond *in time*. *Latency attacks*—also called energy-latency, slowdown, sponge, or resource-consumption attacks—deliberately inflate the inference-time computation, energy, or wall-clock latency of a deep model, so that a real-time consumer (a vehicle controller, an interactive service, a battery-powered sensor) misses its deadline or exhausts its resources even when the underlying predictions remain nominally accurate. This type of attack is especially problematic for machine-learning systems running on embedded systems with limited computational, memory, and battery resources, already under tight deadlines. This survey synthesizes a literature that has, until now, been fragmented across distinct communities: physical-world attacks on autonomous-driving perception, sponge examples on convolutional accelerators, slowdown attacks on input-adaptive (multi-exit) networks, efficiency-degradation attacks on neural machine translation and image captioning, and—most recently—output-length inflation, verbose-image, and “infinite-thinking” denial-of-service attacks against large language models (LLMs), vision-language models (VLMs), mixture-of-experts (MoE) routers, and chain-of-thought reasoning models. We organize this body of work along two orthogonal axes: the *exploited bottleneck* (post-processing, temporal state, sensor/cooperative fusion, hardware, dynamic-computation control flow, autoregressive generation, and training-time poisoning) and the *model class* affected (CNN detectors, dynamic/adaptive networks, transformers, LLMs, and multimodal generators). We present a unified threat-model taxonomy, consolidated quantitative comparisons of reported latency, energy, and amplification factors, and a structured review of the nascent defense landscape—budget enforcement, robust/adaptive training, input transformation, architectural sparsity, runtime monitoring, and system-level serving control. We close with a discussion of cross-cutting open challenges: the absence of standardized benchmarks reporting both timing *and* downstream safety impact, the gap between digital and physically realizable attacks, the new attack surface created by inference-time compute scaling, and the tension between efficiency and robustness that pervades every proposed defense. A companion website is available here: <https://github.com/guzonghua/awesome-latency-attacks/>.

Additional Key Words and Phrases: availability attacks, energy-latency attacks, sponge examples, inference latency, denial-of-service, object detection, non-maximum suppression, multi-exit networks, large language models, vision-language models, mixture of experts, real-time systems, adversarial machine learning

1 Introduction

Deep neural networks are now deployed in settings where *when* a prediction arrives is as important as *whether* it is correct. An autonomous vehicle that perceives a pedestrian one second late has effectively not perceived it at all; an interactive assistant that takes thirty seconds to begin responding is unusable; a battery-powered sensor whose model suddenly draws ten times its budgeted energy will exhaust itself before its mission completes. These deployments expose an attack surface that classical adversarial machine learning largely ignored.

Most adversarial-example research targets *integrity*: an attacker perturbs an input so that the model misclassifies it [84]. *Availability* attacks pursue a different goal. Rather than changing the answer, the attacker inflates the cost of producing it—measured in floating-point operations, energy, or latency—so the system cannot respond “within a reasonable time” [7]. We use *latency attack* as an

umbrella term for this family, which also appears in the literature under the names energy-latency attack, sponge attack, slowdown attack, efficiency-degradation attack, and resource-consumption denial-of-service. The unifying mechanism is the deliberate steering of a model, or its host system, toward worst-case computational behavior.

The overall organization of this survey is summarized in Figure 1.




Fig. 1. Overview of the survey structure.

1.1 Why latency attacks matter now

Although many works surveyed in this paper are recent, the rapid expansion of latency attacks across perception, adaptive inference, and foundation models motivates a unified survey at this stage. Three trends have converged to make latency attacks a first-class concern. First, *real-time autonomy* couples perception latency directly to physical safety: a delayed detection cascades into delayed planning and control [60]. Second, the field has embraced *input-adaptive efficiency*—early-exit networks, token pruning, skimming transformers, and mixture-of-experts routing—which trade a fixed worst-case cost for a lower *average* cost. This is precisely the property an availability attacker can revoke: by forcing every input onto the slow path, the attacker erases the advertised savings and may even exceed the cost of the static model [12, 34]. Third, the rise of *autoregressive generation* in LLMs and VLMs means that a small input can induce an arbitrarily large output, and the per-query compute cost is now an attacker-controllable variable [16, 25]. The recent move to inference-time “thinking” in reasoning models multiplies this leverage: a single innocuous-looking prompt can trigger tens of thousands of reasoning tokens [55].

1.2 Scope and contributions

Prior surveys treat pieces of this landscape in isolation. The closest, by Brachemi Meftah et al. [5], covers energy-latency attacks on regular, input-adaptive, and decoder-based models, but predates and only shallowly covers the LLM/VLM era—reasoning-model denial-of-service, mixture-of-experts routing attacks, and multimodal verbose attacks. The autonomous-driving security literature, meanwhile, has produced a rich set of *physically realizable* latency attacks that are rarely discussed alongside the general deep-learning work. This survey deliberately bridges both. We make the following contributions:

- **A unified taxonomy** that organizes latency attacks by exploited bottleneck and by affected model class, spanning object detection, dynamic networks, transformers, LLMs, and multimodal generators.
- **A unifying principle**—*intermediate-work amplification* (Section 2.3)—that explains every attack class through a single mechanism and motivates a cross-domain *work-budget* defense abstraction.
- **A threat-model framework** (Section 3) that classifies the attacker interface, knowledge, and targeted pipeline stage, making otherwise incomparable works commensurable.
- **Consolidated quantitative comparisons** (Tables 3, 4, 5) of reported latency, energy, FLOP, and output-length amplification factors, with explicit notes on metric incompatibility.
- **A structured defense review** (Section 8) organized by control mechanism rather than by target, exposing which point of the compute pathway each defense actually bounds.
- **A cross-domain synthesis and evaluation discussion** (Section 9) showing why perception-side and serving-side defenses do not transfer, including a review of fragmented evaluation practice (Section 9.1).
- **A forward-looking analysis** (Section 10, including the NMS-free detector discussion in Section 10.1) of emerging attack surfaces and the open challenges that define the next phase of this research.

2 Background and Definitions

2.1 Integrity versus availability

Deep learning is now pervasively deployed on resource-constrained embedded and edge platforms, where inference latency, energy, and throughput are first-class design constraints rather than afterthoughts [44, 57, 62]. It is precisely in these latency- and energy-sensitive regimes that availability attacks become consequential. We adopt the NIST taxonomy of adversarial machine learning, which distinguishes integrity, availability, and confidentiality violations [84]. Integrity attacks aim to change the model’s output; availability attacks aim to degrade its usefulness or serviceability. Latency attacks are availability attacks whose specific lever is computational cost. Crucially, a latency attack may leave the prediction *unchanged*—this is a feature, not a side effect, because output-quality monitors (BLEU for translation, mAP for detection, perplexity for text) will not flag an attack that only changes timing [10, 75]. Gao et al. [28] evaluated and compared well-known certifiable defenses against adversarial patches across different hardware platforms with varying computing power in terms of both robust accuracy and processing time of the defense algorithm, which addresses availability attacks and is orthogonal to the availability attacks considered in this survey article.

2.2 Where latency comes from

A modern inference pipeline contains many cost centers, and latency attacks succeed by driving one or more toward its worst case:

Data-dependent post-processing. Operations whose cost depends on the *content*, not just the shape, of intermediate tensors. The canonical example is non-maximum suppression (NMS) in object detectors, whose cost grows with the number of candidate boxes [75].

Activation sparsity. Hardware accelerators (GPUs, TPUs, ASICs) exploit zero-valued activations to skip computation, and approximate-computing techniques trade accuracy for energy on the edge [29]. Inputs that minimize sparsity force maximum multiply-accumulate work and energy [64, 76].

Conditional computation. Early-exit networks, token-pruning vision transformers, skimming language models, and mixture-of-experts routers all decide *at runtime* how much computation to perform [19, 70, 82]; the same input-adaptive model-selection strategies that save average-case energy on embedded systems [62] create the worst-case headroom an attacker can reclaim. Each decision point is an attack target.

Autoregressive generation. Sequence models emit one token at a time until an end-of-sequence (EOS) signal. The number of tokens—and thus total cost—is determined by the model’s own outputs, which the attacker can steer [10, 25].

Temporal and multi-agent state. Pipelines that maintain state across frames (multi-object tracking) or aggregate across agents (cooperative perception) add stages whose cost depends on the number of tracked or fused entities [40, 86].

Hardware and memory. Memory access patterns, DRAM bandwidth, and even induced bit-flips can be manipulated to inflate latency below the algorithmic level [66, 77].

2.3 The unifying principle: intermediate-work amplification

Beneath the diversity of targets and modalities lies a single mechanism, which we call *intermediate-work amplification*. In every latency attack, the adversary engineers a situation in which some internal boundary of the computation receives more objects, tokens, attention computations, or reasoning steps than a benign input of the same nominal class would generate. In object detection, those objects are phantom bounding boxes that flood NMS or multi-object tracking. In transformer inference, they are tokens that occupy the KV-cache and drive quadratic attention. In autoregressive decoding, they are extra generated tokens; in reasoning models, extra chain-of-thought steps. The attack is effective precisely because the relevant algorithms—NMS, self-attention, autoregressive decoding, Hungarian matching—have at least *super-linear worst-case complexity* in the count of intermediate objects, so a modest increase in the count produces a disproportionate increase in cost. This shared structure appears across both perception pipelines and LLM/VLM-serving domains. This abstraction is more than descriptive: it implies that the natural cross-domain defense is a *work budget*—an explicit cap on the number of intermediate objects, tokens, or attention computations permitted per unit time, enforced at the boundary between stages (Section 9).




Fig. 2. Overview of latency attacks as system-level availability threats.

Figure 2 presents an overview of latency attacks as system-level availability threats. Deployed AI systems operate under latency, energy, throughput, and deadline budgets. Attackers exploit data-dependent bottlenecks to amplify intermediate work and push execution toward high-cost paths, causing latency spikes, energy drain, throughput collapse, missed deadlines, service degradation, or timing leakage. Defenses restore bounded and predictable computation through budget enforcement, monitoring, scheduling, and bounded execution.

2.4 Related surveys

Adversarial-ML surveys overwhelmingly emphasize *evasion*—perturbations that cause misclassification—and treat the availability axis as peripheral [1, 30, 52]. Surveys of adversarial attacks on vision-language models catalog jailbreak, backdoor, and patch threats but largely omit output-length and resource-consumption attacks [14, 23, 35]. LLM-security surveys catalog serving-layer availability attacks but do not connect them to the perception-stack latency literature [99]. Earlier work by Kianpour and Wen [43] surveyed *timing attacks* on ML systems and laid conceptual groundwork, but predates the NMS-overload and LLM-DoS literature. The closest dedicated treatment, Brachemi Meftah et al. [5], covers energy-latency attacks on regular, adaptive, and decoder-based models but only addresses the LLM/VLM era in a shallow way. To our knowledge, this survey is the first to unify latency attacks and defenses across perception pipelines, dynamic networks, and LLM/VLM serving under a single threat model. Table 1 positions our work against prior surveys along three coverage axes—autonomous-driving (AD) perception, general deep learning, and the LLM/VLM era—and identifies the principal gap each leaves open.

Table 1. Positioning of this survey against prior surveys touching latency, timing, or availability attacks. ●: substantial coverage; ◐: partial/peripheral coverage; ○: not covered.

Survey	Year	AD percep.	General DL	LLM/VLM	Principal gap
Kianpour & Wen [43]	2019	○	◐	○	Predates modern latency/sponge and LLM-DoS literature
NIST AML taxonomy [84]	2023	○	◐	◐	Definitional only; availability is one small category
LLM-security surveys [1, 30, 52, 99]	2024–25	○	○	●	Serving-layer DoS only; no link to perception latency
VLM adversarial surveys [14, 23, 35]	2025–26	◐	○	◐	Jailbreak/backdoor/patch focus; omits output-length & resource attacks
Brachemi Meftah et al. [5]	2025	◐	●	◐	Shallow on LLM/VLM, reasoning, and MoE; no AD system-level synthesis
This survey	2026	●	●	●	— (unifies all three via intermediate-work amplification)

2.5 A note on metrics

Reported results in this field are often incomparable. Authors may report multiplicative latency increases ($\times$), absolute milliseconds, FLOP increases, energy in joules, generated-token counts, frames per second, and downstream task outcomes (crash rate, collision rate). On embedded targets, the relevant quantity is often the worst-case execution time relative to a real-time deadline, which has motivated dedicated execution-time prediction frameworks for concurrent DNN workloads on edge accelerators [31]. Throughout, we preserve each paper’s native metric and flag where direct comparison is unsound. Establishing common axes is itself an open problem (Section 10).

3 Threat Models

A coherent threat model is the prerequisite for comparing attacks that span sensor inputs, public APIs, and hardware faults; without it, headline slowdown numbers are not commensurable. We classify latency attacks along three dimensions: the *attacker interface* (how the adversary touches the system), the *attacker knowledge* (white-, grey-, or black-box), and the *targeted stage*. Table 2 summarizes representative families.

Table 2. Threat-model families for latency attacks, with representative works.

Family	Interface	Targeted stage	Representative evidence
Digital input perturbation	Modify sensor input (image, point cloud, text)	Detector inference / NMS / decoder	Overload [7]; Phantom Sponges [75]; sponge examples [76]
Pipeline-level perturbation	Modify camera input for end-to-end effect	Full perception (detect + track)	SlowTrack [60]
Physical projector / patch	Project or print patterns into the scene	Camera perception via phantom objects	SlowPerception [61]; DetStorm [65]
Sensor runtime maximization	Perturb LiDAR point clouds	LiDAR detection runtime	SlowLiDAR [53]
Cooperative V2V injection	Inject perturbations via V2V messages	Cooperative-perception fusion	CP-FREEZER [86]
Autoregressive inflation	Craft input prompt / image	Decoder generation length	NMTSloth [10]; Engorgio [16]; Verbose Images [25]
Dynamic-control defeat	Craft input to defeat conditional compute	Early-exit / pruning / routing	DeepSloth [34]; No-Skim [102]; Misrouter [20]
Training-time poisoning	Poison training data or weights	Persistent inference cost	Sponge poisoning [13]; Skip-Sponge [81]; DoS poisoning [26]
Hardware / side-channel	Induce faults or observe timing	NMS / memory / timing oracle	Bit-flip attack [77]; timing side-channel [66]

3.1 Attacker interface

The interface ranges from purely digital (modifying the input tensor of a model under white-box control) to physical (projecting or printing adversarial patterns that a camera then captures) to

systemic (injecting messages over a vehicle-to-vehicle channel, or submitting prompts to a public API). Physical realizability is the dividing line between laboratory demonstrations and deployable threats, and is a recurring theme in the autonomous-driving literature [61, 65].

3.2 Attacker knowledge

White-box attacks assume gradient access and dominate early work because the optimization is direct. Grey-box attacks assume architectural knowledge but not exact weights [17]. Black-box attacks, increasingly important for deployed services, rely on transfer from a surrogate, on query feedback (e.g., `usage.prompt_tokens`), or even on timing itself as an oracle [66, 105]. The existence of timing-only black-box attacks is significant: it means even an opaque service that returns no probabilities can leak enough information to be optimized against.

4 Latency Attacks on Object Detection and Perception

Object detection is the most thoroughly studied target, because its post-processing stage offers an unusually clean computational bottleneck and because autonomous driving supplies a high-stakes deployment context.

4.1 Non-maximum suppression as a bottleneck

NMS filters overlapping candidate boxes and has a cost that grows with the number of candidates. *Phantom Sponges* [75] crafts a universal adversarial perturbation that adds many “phantom” boxes to overload NMS while *preserving* the original detections, so accuracy collapse does not betray the attack. *Overload* [7] generalizes this by formulating latency maximization as an optimization with a spatial-attention mechanism, reporting that single-image inference time can be increased tenfold, and emphasizing that the attack is “NMS-agnostic.” *Beyond PhantomSponges* [74] modifies a box-area loss to lower intersection-over-union and exacerbate NMS load, reporting up to a 550% increase in NMS time on a YOLOv5-small configuration. *Groundswell* [95] sharpens the stealth–effectiveness tradeoff with a *Regional Perturbation Balance* strategy that injects phantom objects into spatial regions that minimally disturb the original detections; on YOLOv5s it inflates the candidate count fed to NMS from 101 to roughly 19,000 and end-to-end latency from 15.2 ms to 41.4 ms (about a 172% increase), while cutting optimization time by 80% relative to prior universal-perturbation attacks. *Daedalus* [87], although framed as an integrity attack, is a relevant NMS failure mode: it compresses box dimensions so NMS cannot filter dense false positives, and is demonstrated via printed posters. A complementary observation by Biton et al. [4] reframes variable-time inference as a security liability in itself: because NMS runtime is data-dependent, it *leaks* internal inference state, and this timing signal can be used to enhance black-box, decision-based attacks. This dual role—NMS as both a target of latency inflation and a side-channel oracle—motivates constant-time NMS as an explicit security requirement (Section 8).

4.2 Whole-pipeline and temporal-state attacks

Detector-only attacks often fail to change downstream behavior because multi-object tracking (MOT) absorbs detection noise; prior work shows that over 98% detection-attack success may be required to alter tracking outcomes [40]. This motivates pipeline-level threat models. *SlowTrack* [60] attacks camera-based perception end to end with a two-stage strategy and three new loss designs, reporting average latency $2.9\times$ higher than prior approaches for latency of the whole camera-based AD perception system; and, critically, connecting latency to safety: a vehicle crash rate of $\sim 95\%$ versus $\sim 30\%$ for baselines, evaluated on the Baidu Apollo stack and LGSVL simulator. Related temporal-integrity work—tracker hijacking [40, 59] and trajectory-aware attacks on MOT [36]—reinforces that stateful components can amplify or attenuate component-level perturbations.

4.3 Physical-world latency attacks

SlowPerception [61] introduces projector-based universal perturbations that create numerous phantom objects, inducing an average physical-world latency of 2.5 seconds and a 97% collision rate in an industry-grade system plus simulator. *DetStorm* [65] advances physical realizability, optimizing objects against multiple NMS variants and reporting a 506% average increase in detected objects and delays up to 8.1 seconds, while arguing that prior digital or large-patch attacks lack real-world applicability. *Groundswell* [95] is explicitly designed for physical delivery: the universal perturbation can be affixed as a sticker on the camera lens or projected onto a surface ahead of the vehicle, injecting itself into raw frames before in-vehicle processing and pushing per-frame latency past the 30 ms real-time budget required by Level-4/5 autonomy.

4.4 Beyond cameras: LiDAR and cooperative perception

SlowLiDAR [53] presents the first systematic availability attack on LiDAR detection, handling non-differentiable preprocessing via differentiable proxies and an execution-time-aware loss to inflate runtime across six popular pipelines while maintaining imperceptibility. *CP-FREEZER* [86] targets vehicular cooperative perception, injecting adversarial perturbations via V2V messages to maximize fusion delay; it resolves point-cloud nondifferentiability and transmission asynchrony, increasing end-to-end latency by over 90×, pushing per-frame time beyond 3 seconds, and achieving 100% success on a real vehicle testbed. A simulation study further analyzes the system-level impact of perception inference-time attacks on full autonomous vehicle stacks [8], and real-time defenses against object-based LiDAR attacks are beginning to appear [103].

Table 3. Representative latency attacks on object detection and perception. Metrics are reported as in the source; “–” indicates not reported in a directly comparable form.

Method	Target stage	Key quantitative result	Physical?
Phantom Sponges [75]	NMS (end-to-end OD)	Universal perturbation overloads NMS while preserving detections	No
Overload [7]	OD inference (edge)	Single-image inference time up to 10×; NMS-agnostic	Edge-eval
Beyond Phantom-Sponges [74]	NMS runtime (YOLO)	Up to 550% NMS-time increase (YOLOv5-small)	No
Groundswell [95]	NMS (camera OD)	Candidates 101→19k; 15.2→41.4 ms (~172%); –80% opt. time	Yes (sticker/projector)
SlowTrack [60]	Camera perception (detect+track)	Latency 2.9× vs. prior; crash rate ~95% vs. ~30%	Patch (prelim.)
SlowPerception [61]	NMS + MOT workload	Avg. 2.5 s physical latency; 97% collision rate	Yes (projector)
DetStorm [65]	Camera perception delay	506% more detections; up to 8.1 s delay	Yes
SlowLiDAR [53]	LiDAR detection runtime	Significant latency rise on six LiDAR pipelines	No
CP-FREEZER [86]	Cooperative fusion	> 90× latency; >3 s/frame; 100% on testbed	Yes (testbed)
Sponge Backdoor [96]	OD latency via NMS	Poisoned inputs prolong time; clean accuracy preserved	No
Bit-flip attack [77]	NMS via parameters	Up to 71.6 ms (20.4×) with 31 bit-flips	Partial

4.5 Embodied and Vision-Language-Action Availability Attacks

The latest frontier moves past the perception front-end to the *Vision-Language-Action* (VLA) models that map multimodal observations directly to robot actions, where an availability failure is a physical failure: a robot that stalls or acts late can collide, drop a payload, or freeze mid-manipulation. *FreezeVLA* [90] formalizes *action-freezing* attacks via min-max bi-level optimization: a single adversarial image “freezes” the VLA and makes it ignore subsequent instructions—disconnecting the robot’s digital mind from its physical actions—attaining a 76.2% average freeze rate across three VLA models and four benchmarks, with strong cross-prompt transferability. *Semantic-DoS* [79] turns the model’s own safety alignment into the attack surface: injecting short (1–5 token) safety-plausible phrases into the robot’s audio channel triggers protective halting without any jailbreak, reaching up to 98.3% hard-stop success and revealing that prompt-level defenses merely reshape—rather than remove—the disruption. Patch-based variants such as the Embedding Disruption Patch Attack (EDPA) [98] place a model-agnostic patch in the camera view that breaks visual–textual alignment and drives repeated incorrect actions to task failure, and *MAVLA* [106] breaks cross-modal alignment through a coordinated multimodal adversarial framework. *ANNIE* [39] provides the first systematic, ISO-grounded study of embodied safety violations, decomposing long-horizon goals into frame-level perturbations and exceeding 50% success across critical, dangerous, and risky categories with validation on a physical robot. These attacks share the availability signature of the rest of this survey—worst-case work or worst-case delay imposed on a real-time loop—but relocate the consequence from a slow inference to an unsafe or absent physical action; broader treatments of the embodied threat surface are emerging in dedicated surveys [97].

Section Summary. Collectively, these studies illustrate a clear evolution of latency attacks in perception systems. Early work primarily exploited a single computational bottleneck, most notably non-maximum suppression (NMS). More recent research increasingly targets complete perception pipelines, incorporates physical realizability, and evaluates downstream consequences such as tracking failures, delayed planning, and collision risk—and, most recently, extends to the embodied VLA policies that consume perception output, where induced latency or action-freezing becomes a direct physical hazard. This progression suggests that latency attacks are evolving from component-level vulnerabilities into system-level threats that must be evaluated across the entire autonomy stack. More importantly, the object-detection literature also reveals a broader lesson: although NMS happened to be the first successful target, the underlying objective is not to attack NMS itself, but to maximize intermediate computation wherever it occurs. The following sections show that the same principle extends naturally to activation sparsity, adaptive inference, autoregressive generation, and LLM serving.

5 Neural Inference Work-Inflation Attacks

The *sponge* lineage targets the general mechanism of activation sparsity rather than any specific post-processing stage, making it broadly applicable across architectures.

5.1 Inference-time sponge examples

Sponge examples [76] are the seminal work: adversarial inputs crafted to maximize energy and latency by minimizing the activation sparsity that accelerators exploit. On transformer and vision models, they report up to 30× increases in latency and energy consumption in some settings, and a real-world demonstration that raised a cloud translator’s response time from 1 ms to 6 s. A follow-up analysis shows that uniform/constant inputs act as “free” sponge examples and that the achievable energy increase is bounded by the architecture’s sparsity range, motivating activation-sparsity monitoring as a defense signal [64].

5.2 Sponge poisoning: moving the attack to training

Sponge poisoning [13] injects a small fraction of crafted samples so the trained model incurs elevated energy on *all* natural inputs thereafter—a far stealthier and more persistent variant. *SkipSponge* [81] operates directly on weights, poisoning residual (skip) connections with under 1% weight modification and evading fine-pruning and parameter-perturbation defenses. Sponge poisoning has been adapted to on-device DNNs [94], to multi-exit networks (where it both prevents early exits and raises per-layer cost) [37], to mobile apps [69], and to FPGA accelerators under differential privacy [2]. A sensing-AI study quantifies battery-drain risk in IoT settings and proposes pruning as a defense, reducing energy overhead from $\sim 4\times$ to $\sim 1.3\times$ [33].

5.3 Sparsity and spiking networks

Sparsity attacks [45] degrade energy/latency by minimizing activation sparsity and withstand activation thresholding and input quantization. Spiking neural networks, attractive for low-power edge perception, are also vulnerable: efficiency attacks raise spike activity by $1.7\times$ – $2.5\times$ and energy by $1.4\times$ – $2.2\times$ [46], and timestep-compressed methods reduce the latency of *generating* such attacks for real-time use [42].

5.4 Dynamic and Adaptive Networks

Input-adaptive networks promise lower *average* cost; latency attacks revoke that promise by forcing worst-case control flow.

Multi-exit networks. The earliest availability attack on input-adaptive networks is *ILFO* [32], which optimizes an intermediate-output loss to reactivate skipped blocks in depth-adaptive networks, increasing FLOPs by up to 100%. *DeepSloth* [34] is the foundational slowdown attack on early-exit inference: per-input perturbations force the network to defer to its final, most expensive classifier, eliminating the efficiency advantage with 90–100% success and up to $5.2\times$ cost inflation. Refined energy attacks combine energy maximization with confidence minimization and remain effective in grey-box settings [17]. In NLP, *SlowBERT* [101] attacks multi-exit BERT, and *Dynamic Transformers Provide a False Sense of Efficiency* [12] introduces a unified attack (SAME) that defeats early exits across BERT, RoBERTa, ALBERT, and DistilBERT with up to $5\times$ slowdown. A foundational study demystifies *why* these models are vulnerable and proposes calibration-based defenses that recover roughly 60% of lost efficiency [83].

Token pruning, skimming, and mixture-of-experts. Skimming language models drop unimportant tokens; *No-Skim* [102] crafts inputs that prevent dropping, raising FLOPs by 2 – $3\times$. A unified framework treats early exit, token-pruning vision transformers, and MoE routing as instances of one “efficiency attack” class, demonstrating 2 – $8\times$ overhead across all three [71]. In the vision-transformer setting specifically, *SlowFormer* [67] trains a universal adversarial patch (as little as 2% of the image area) that drives efficient ViTs such as A-ViT, ATS, and AdaViT toward their maximum compute, and *QuantAttack* [3] exploits dynamic quantization to force worst-case high-precision execution, inflating inference time, memory, and energy. Mixture-of-experts routers, central to modern sparse LLMs, are a fresh target: *Misrouter* [20] causes routing collapse or saturation in a black-box, query-only setting, yielding 2 – $4\times$ latency increases, while *RouteHijack* [100]—though primarily a jailbreak attack—demonstrates that adversarial control of expert routing can additionally inflate computation by activating costly experts. *RepetitionCurse* [38] shows the vulnerability is systemic rather than incidental: under expert parallelism, simple repetitive-token prompts force all tokens onto the same top- k experts, creating device-level bottlenecks while others idle and raising end-to-end latency by $3.06\times$ on Mixtral-8x7B—turning the efficiency mechanism itself into a

denial-of-service vector. *SPLAT* [48] revisits latency attacks on dynamic networks with a black-box, stealthy formulation for multi-exit models.

Section Summary. Unlike perception attacks, sponge attacks reveal that latency vulnerabilities are not tied to a particular application or post-processing stage. Instead, they exploit architectural properties shared by many neural networks, including activation sparsity and conditional computation. Training-time variants further demonstrate that excessive computational cost can be embedded persistently into model parameters rather than triggered solely by adversarial inputs. Taken together, these studies shift the focus from application-specific bottlenecks toward general properties of efficient neural inference.

6 Autoregressive and Serving-Side Work Inflation

6.1 Sequence Generation and Decoder-Length Inflation

Autoregressive decoders introduce a uniquely powerful lever: the attacker can inflate the *number of generated tokens*, and total cost scales with it.

NMTSloth [10] is the first systematic study of efficiency robustness in neural machine translation, crafting source sentences that inflate decoder steps by 5–10× and end-to-end latency by 3–8× while leaving BLEU nearly unchanged so no quality monitor fires. *LLMEffiChecker* [21] extends this to general LLMs, finding them *more* vulnerable: a five-token perturbation can inflate generation length by 10–30×. *NICGSlowDown* [11] brings the paradigm to neural image captioning, forcing decoders to emit far longer captions via image perturbations, and establishes a widely reused efficiency-robustness benchmark. *TTSslow* [27] demonstrates that the threat reaches audio: adversarial text inflates autoregressive text-to-speech decoding by up to 5×. The common thread—and a recurring defensive blind spot—is that none of these attacks degrade output quality, so they evade quality-based detection entirely.

6.2 LLM and VLM Resource-Consumption Attacks

The LLM/VLM era both inherits the autoregressive lever and adds new ones: chain-of-thought reasoning, multimodal inputs, and shared serving infrastructure. The emergence of reasoning models fundamentally changes the threat landscape. Earlier latency attacks manipulated inference pipelines whose computational cost was largely bounded. Modern reasoning models intentionally expose computation as an input-dependent resource, giving attackers much greater leverage.

Output-length inflation and inference-cost DoS. *Engorgio* [16] optimizes prompts that coerce LLMs into babbling, inflating output token counts by 10–50× and transferring to black-box targets. *Crabs* [105] automates black-box LLM-DoS prompt generation, sustaining 10–30× output inflation even under provider-side length monitoring through semantic camouflage. *Denial-of-Service Poisoning* [26] shows the threat can be embedded at training time: a triggered model emits endless output, never producing EOS, with 100% success at under 1% poisoning and survival through RLHF safety tuning.

Reasoning-model and “infinite-thinking” attacks. Inference-time compute scaling creates a new, high-leverage surface. *ReasoningBomb* [55] induces pathologically long reasoning traces in large reasoning models—averaging ~18,759 completion tokens versus ~300 for normal queries, a roughly 60× cost multiplier—while remaining stealthy and bypassing input-, output-, and joint-detection baselines at rates of 99.8%, 98.7%, and 98.4%, respectively. *ThinkTrap* [49] exploits chain-of-thought mode in black-box APIs to induce near-infinite thinking loops, consuming 50–100× more tokens than standard requests. These results strongly argue for budget enforcement through *construction* rather than detection.

Multimodal verbose attacks. *Verbose Images* [25] crafts adversarial images that make large VLMs (GPT-4V, LLaVA, InstructBLIP, MiniGPT-4) generate long, repetitive responses, inflating output by 10–30× and latency by 7–30× while preserving visual fidelity. *VLMInferSlow* [92] provides a service-oriented evaluation across 12 open and closed VLMs, reporting 85–100% success and 5–40× increases in latency. *Verbose-Text Induction* [58] co-optimizes adversarial images and prompts to exceed 10,000 output tokens from a single query, and *LingoLoop* [22] traps multimodal models in self-referential generation loops that fill the context window faster than legitimate output, defeating context-length caps. The trajectory from image captioning [11] through verbose images to multimodal loop attacks is a coherent and accelerating escalation. This escalation now extends to the temporal domain: *VidDoS* [80] is the first universal energy-latency attack tailored for video-based LLMs, using a single instance-agnostic adversarial patch—trained via masked teacher forcing with a refusal penalty and early-termination suppression—that requires no inference-time gradients and therefore keeps pace with continuous video streams. It induces token expansion exceeding 205× and latency inflation exceeding 15× across three Video-LLMs, and simulations of real-time autonomous-driving streams show the induced latency causes critical safety violations. The threat has also reached *embodied* multimodal autonomy: a natural-reflection backdoor attack on driving VLMs [54] embeds reflection-pattern triggers in camera images that, when present, cause the VLM planner to generate abnormally long responses—demonstrating significantly increased inference latency under trigger with direct real-world safety implications, and uniting the output-length-inflation and training-time-backdoor threat classes in a single AD-relevant attack.

Serving-framework, tool-chain, and agentic DoS. The newest attacks move below the model to the *serving system* and the *agentic loop*, where shared infrastructure amplifies a single adversary’s leverage. *Fill and Squeeze* [88] reports the counter-intuitive finding that continuous batching largely neutralizes purely algorithmic length attacks, then attacks the scheduler directly—exhausting the global KV cache to induce head-of-line blocking and forcing repetitive preemption—for up to 20–280× time-to-first-token slowdown in a black-box setting at lower cost than input-only methods. At the agent layer, *Beyond Max Tokens* [107] exploits the Model Context Protocol tool interface: text-only edits to tool metadata, optimized by Monte-Carlo tree search, steer agents into prolonged tool-calling chains exceeding 60,000 tokens, raising per-query cost by up to 658× while preserving task success and evading prompt filters. *From Shield to Target* [108] turns LLM *guardrails* into the victim: crafted payloads trap the guardrail in extended reasoning loops for 13–63× token amplification and up to 148× latency amplification, and a single poisoned document can starve co-located agents sharing the guardrail. *Mobius Injection* [51] generalizes this to agent-based DDoS (AbO-DDoS): by exploiting a “semantic closure” vulnerability, one textual injection induces sustained recursive agent execution, driving single-node call amplification up to 51× and multi-node p95 latency up to 229×. These developments echo the broader systematization of the agentic attack surface [15] and are now recognized at the standards level as “unbounded consumption” [68]—confirming that availability, not just integrity, is a first-class security property of deployed LLM services.

Section Summary. Earlier attacks were constrained by the computational graph of a fixed model, whereas modern sequence models expose output length—and increasingly reasoning length—as attacker-controllable variables. As a result, computation itself becomes part of the attack surface. This transition explains why resource-consumption attacks have rapidly expanded from machine translation and image captioning to LLMs, VLMs, reasoning models, and mixture-of-experts architectures.

Table 4. Representative latency/energy attacks beyond object detection. Native metrics preserved.

Method	Domain / target	Key quantitative result
Sponge Examples [76]	General DNN (vision, NLP)	Up to 30× energy/latency; cloud translator 1 ms→6 s
Sponge Poisoning [13]	Training-time, general	1–5% poison ⇒ persistent 2–10× energy
SkipSponge [81]	Weight poisoning	High energy overhead at <1% weight change; evades defenses
DeepSloth [34]	Multi-exit networks	90–100% force full inference; up to 5.2× cost
Dynamic Transformers (SAME) [12]	Multi-exit NLP	Up to 5× slowdown across BERT family
No-Skim [102]	Skimming LMs	2–3× FLOPs; output label unchanged
Misrouter [20]	MoE routing (black-box)	2–4× latency via routing collapse/saturation
RepetitionCurse [38]	MoE expert parallelism	3.06× latency on Mixtral-8x7B via router imbalance
NMTSloth [10]	Neural machine translation	5–10× decoder steps; 3–8× latency; BLEU stable
NICGSlowDown [11]	Image captioning	5–10× decoding steps; 3–7× latency
Engorgio [16]	LLM prompts	10–50× output tokens; black-box transfer
Crabs [105]	LLM-DoS (black-box)	10–30× output; evades length monitoring
ReasoningBomb [55]	Reasoning models	~60× tokens; >98% detection bypass
Verbose Images [25]	VLMs	10–30× output; 7–30× latency
VLMInferSlow [92]	VLM-as-a-service	85–100% success; 5–40× latency
Fill and Squeeze [88]	LLM serving scheduler	20–280× TTFT slowdown; black-box
Beyond Max Tokens [107]	LLM agent tool chains (MCP)	Up to 658× per-query cost; >60k-token chains
From Shield to Target [108]	LLM guardrails	13–63× tokens; up to 148× latency
Mobius Injection [51]	Agentic AbO-DDoS	51× call amplification; 229× p95 latency

7 Training-Time Persistence, Hardware, and Timing Channels

Not all latency attacks act on the test-time input. The attacks discussed in this section broaden the threat model beyond the test-time input: they persist across sessions, operate below the algorithm, or turn latency itself into an information leak, and are therefore the hardest to detect and remove. *Training-time* variants (sponge poisoning [13], SkipSponge [81], sponge backdoors in detectors [96], and no-EOS DoS poisoning of LLMs [26]) introduce persistent cost inflation that persists through normal usage and is hard to remove without retraining. *Hardware* attacks operate below the algorithm: a row-hammer bit-flip approach targets NMS parameters via side channels, reporting up to 71.6 ms (20.4×) latency with just 31 bit-flips [77], and FPGA sponge poisoning evades differential-privacy noise to inflate accelerator energy [2]. *Side-channel* attacks invert the relationship between timing and computation: a timing-based black-box method crafts adversarial examples using only

processing time, exploiting the link between activated nodes and latency, and evades gradient-masking countermeasures [66]. This last result is conceptually important—it shows that latency is not only a target but also a leak.

8 Defenses

Defenses are nascent, often attack-family-specific, and pervaded by an efficiency–robustness tradeoff. We organize them by *control mechanism*—the point in the compute pathway each one actually bounds—rather than by the attack they were designed to counter. Table 5 summarizes the landscape.

8.1 Budget enforcement (bound the output / compute directly)

The most direct defense against autoregressive inflation and reasoning DoS is to enforce a budget by construction. *PD3F* [104] is a two-stage, pluggable framework that combines input-side dynamic request scheduling (via a “Resource Index”) with output-side “Adaptive End-Based Suppression”, which terminates excessive generation early, improving access capacity by up to 500% under adversarial load across six models. *Token-budget-aware pool routing* [9] estimates each request’s token budget and routes it to a right-sized pool, avoiding worst-case provisioning that wastes 4–8× concurrency. *Conformal Thinking* [91] reframes the reasoning budget as risk control, stopping when the model is confident or when an instance is likely unsolvable, directly countering “runaway” reasoning. Because detection-based defenses are demonstrably fragile—ReasoningBomb bypasses them at >98% [55]—budget enforcement by construction is the more robust posture.

Among existing defenses, budget enforcement appears to be the most direct approach to bounding worst-case computation, rather than merely reducing average computational cost. We therefore expect it to become increasingly important as reasoning models continue to expose input-dependent computation.

8.2 Robust and adaptive training

Hardware-aware adversarial training [89] builds background attention into the training pipeline using objectness loss as a proxy, restoring real-time throughput from 13 FPS to 43 FPS under attack on a Jetson Orin NX, with a better clean/robust tradeoff. For multi-exit models, *calibration-based defenses* [83] (temperature scaling, robust exit calibration) recover roughly 60% of the efficiency lost to slowdown attacks.

8.3 Input transformation and purification

For NMS-exploitation attacks, *adaptive input resizing* [74] reduces NMS execution time by 90.18% on an attacked YOLOv5-large model. For physical patch/projector attacks, real-time defenses localize and remove adversarial patches via temporal consistency [63] or scale-adaptive DCT frequency analysis with inpainting [72], and integrated attack/defense studies report large mAP recovery and reduced detection time via texture-feature denoising [50]. For VLM visual encoders, *latent-diffusion purification* [47] projects inputs into a low-dimensional latent space to suppress off-manifold perturbations while preserving semantics.

8.4 Bounded and constant-time execution

A first-principles defense is to remove the cost’s data-dependence entirely. Because many bottleneck operations have data-dependent runtime, rewriting them with a fixed worst-case bound eliminates both the latency-inflation lever and the timing side-channel. Biton et al. [4] explore constant-time NMS to close the timing-leakage channel they identify. More generally, capping the number of detections admitted to NMS, or the number of tokens a decoder may emit, converts an unbounded

worst case into a bounded one—an instance of the *work budget* abstraction (Section 9). The cost is a potential accuracy ceiling in dense scenes, making the bound a tunable safety/utility parameter.

8.5 Architectural sparsity and efficiency reduction

Shrinking per-step compute reduces an attacker’s headroom. *Model pruning* [33] improves resilience to sponge poisoning (energy overhead $\sim 4\times \rightarrow \sim 1.3\times$) at an accuracy cost. *BlindSight* [78] exploits attention sparsity to cut VLM FLOPs by 32–41% with -2% to $+2\%$ accuracy. These reduce absolute latency impact but do not *bound* generation length, so they complement rather than replace budget enforcement.

8.6 Runtime monitoring and adaptive inference

Activation-sparsity monitoring offers a lightweight, architecture-agnostic anomaly signal for sponge inputs [64]. For early-exit systems, online scheduling treats the exit decision as a contextual bandit [41] (improving overall performance by up to 98.1% over the best static benchmark), and quorum-based ensemble stopping (SQUAD) [24] cuts inference latency by up to 70.6% while improving accuracy, by avoiding reliance on individual confidence scores that attackers manipulate.

8.7 System-level and serving control

At the serving layer, request scheduling, admission control, and pool right-sizing [9, 104] prevent operational bottlenecks from becoming attack amplifiers. A cautionary result: *certifiers*—formal robustness verifiers with abstain/fallback behavior—paradoxically create a new availability surface, since an adversary can systematically trigger abstention to deny service [56]. Any safety fallback is itself a target. Beyond admission control, hardware-level system defenses target the serving boundary directly: *Garrison* [93] is a GPU-accelerated inference system for adversarial ensemble defense that uses Multi-Instance GPUs with reinforcement-learning-based scheduling, improving adversarial robustness by up to 24.5% while accelerating ensemble inference by $6.6\times$ —a rare example of a defense that improves robustness *and* latency simultaneously. *PSML* [73] secures model-serving systems against model-extraction via fingerprinting and noise-based defenses, providing a complementary protection layer at the inference-serving boundary. The *time-traveling* defense [18] is an unusual input-side variant for the perception domain: it queries historical Street View imagery; it uses majority-vote inference to resist adversarial manipulation of traffic signs, achieving 100% defense effectiveness in its evaluation, demonstrating that temporal redundancy can substitute for per-frame robustness in geographically static scenes.

Section Summary. Existing defenses differ considerably in where they intervene along the inference pipeline, but none provides comprehensive protection. Most approaches either reduce average computational cost or detect anomalous behavior without explicitly bounding worst-case execution. In contrast, mechanisms that directly constrain computation—such as token budgets, bounded execution, or constant-time implementations—appear more promising because they address the underlying cause of latency inflation rather than its symptoms. Nevertheless, balancing efficiency, robustness, and predictive performance remains an open challenge across all defense families.

9 Cross-Domain Synthesis and Evaluation

Viewing the perception and LLM/VLM-serving literatures side by side yields three observations that neither community arrives at independently.

Perception defenses do not transfer to serving, and vice versa. Hardware-aware adversarial training [89], input purification [47, 72], and the time-traveling majority vote [18] all assume the attack

Table 5. Defense families against latency/energy attacks, by control mechanism.

Family	Mechanism	Representative work & signal
Budget enforcement	Cap / terminate generation; route by budget; risk-control stop	PD3F (up to 500% capacity) [104]; token-budget routing [9]; Conformal Thinking [91]
Robust / adaptive training	Background-attentive & hardware-aware AT; exit calibration	Can't Slow Me Down (13→43 FPS) [89]; calibration (~60% recovery) [83]
Input transformation	Resize; patch localization; latent purification; temporal redundancy	Resizing (~90% NMS time) [74]; ADAV [63]; DCT [72]; LDP [47]; time-traveling (100%) [18]
Bounded execution	Constant-time / capped post-processing	Constant-time NMS [4]; detection / token caps
Architectural sparsity	Pruning / quantization; attention sparsity	Pruning ($4 \times \rightarrow 1.3 \times$) [33]; BlindSight (−32–41% FLOPs) [78]
Runtime monitoring	Sparsity/latency anomaly; online exit scheduling	Sparsity monitor [64]; DEE [41]; SQUAD (−70.6% latency) [24]
System / serving control	Scheduling; admission; GPU ensemble; certifier hardening	PD3F [104]; routing [9]; Garrison (+24.5% robust, 6.6×) [93]; PSML [73]; certifier caveat [56]

arrives through a *sensor input* as a pixel-level perturbation; they have no lever on prompt-injection DoS or output-length inflation, where the malicious content is carried in the semantics of a query rather than in pixel noise. Conversely, PD3F [104], token-budget routing [9], and conformal stopping [91] assume a serving stack with request queues and preemptible generation, which has no analog in a camera-to-detector pipeline. This bidirectional non-transfer is the central practical obstacle to securing a heterogeneous autonomy stack.

The intermediate-work amplification principle is universal, and so is its remedy. As argued in Section 2.3, every attack class derives its leverage from forcing a downstream stage to receive more intermediate objects than it was provisioned to process. The natural cross-domain abstraction is therefore a *work budget*: an explicit limit on the number of intermediate objects admitted to stages whose computational cost grows super-linearly with that count.

We view the work budget as an abstract design principle rather than a concrete algorithm. Existing mechanisms—including constant-time NMS [4], detection caps, token budgets [9], and conformal reasoning bounds [91]—can all be interpreted through this common perspective. Although their implementations differ, they share the same objective: explicitly bounding worst-case computation rather than relying solely on average-case efficiency. Viewed in this way, latency defenses become less about protecting individual models and more about controlling computation throughout the inference pipeline. Figure 3 illustrates the work-budget abstraction with perception and serving examples.




Fig. 3. Conceptual illustration of the work-budget abstraction across perception and serving systems. Across both perception and serving systems, limiting the number of intermediate objects admitted to computationally expensive stages prevents adversarial inputs from driving unbounded computation growth. The left panel illustrates candidate-box capping for NMS (e.g., Groundswell on YOLOv5s [95]), while the right panel illustrates token-budget enforcement during autoregressive decoding (e.g., ReasoningBomb [55]). Both panels use log-log axes; the capped curve (dashed) plateaus at the vertical cap line while the unbounded curve (solid) keeps rising. The figure is conceptual and intended to illustrate the common design principle underlying several existing defenses.

No end-to-end defense exists. Every documented defense (Table 5) protects one layer and leaves others exposed. A VLM-equipped autonomous vehicle faces simultaneous latency threats at the camera-to-detector boundary (NMS overload, physical projectors), the tracking layer (SlowTrack-style temporal attacks), the LiDAR and cooperative-fusion stages, and the VLM planner’s decoder (verbose images, reflection backdoors). Composing defenses across these layers is technically possible but has not been evaluated end-to-end in any published work—a gap we return to below.

Takeaway. Viewed through this lens, latency attacks are less about individual algorithms than about uncontrolled computational growth. Although different application domains expose different bottlenecks, they repeatedly exhibit the same underlying mechanism. This shared perspective motivates the work-budget abstraction and suggests that future defenses may be designed around controlling computation itself rather than protecting individual models. We believe this perspective shifts the focus of future research from defending individual models toward bounding computation throughout the entire inference pipeline.

9.1 Evaluation Practices and Benchmarks

Progress is hampered by fragmented evaluation. Perception-attack papers use datasets such as MOT17, nuScenes, and KITTI, report wall-clock latency on specific GPUs (frequently an NVIDIA 2080 Ti or a Jetson device), and—in the strongest cases—report system-level outcomes on the Baidu Apollo stack and LGSVL simulator [60]. LLM/VLM-serving evaluations instead use vLLM-style serving benchmarks on models ranging up to Llama-3-70B, and report access capacity, token amplification, or detector-bypass rate [55, 104]. Three deficiencies recur. First, there is *no standardized latency threat model* analogous to the ℓ_p -ball used for evasion: papers leave attacker access, perturbation/query budget, and the target latency metric unspecified or mutually inconsistent. Second, there is *no agreed success metric*—wall-clock time, FPS, FLOPs, energy, crash rate, and access capacity

are used interchangeably, defeating cross-paper comparison. Third, *no benchmark jointly stresses perception and serving* under attack, which is precisely what an end-to-end AD-plus-VLM stack would require. We argue that joint latency/robustness protocols—measuring accuracy and timing *simultaneously* under adversarial load—are a prerequisite for trustworthy comparison, because efficiency and robustness interact non-monotonically (Section 10).

10 Emerging Surfaces and Open Challenges

10.1 NMS-Free Detectors and Moving Bottlenecks

A salient shift threatens to invalidate much of the object-detection attack literature: modern detectors increasingly remove NMS. YOLOv10 [85] and DETR-style models [6] replace NMS with one-to-one matching (Hungarian or dual-label assignment) and lightweight decoder heads, eliminating the quadratic post-processing blowup that Phantom Sponges and Overload exploit. This does not eliminate the threat—it relocates it. Plausible new bottlenecks include: the multi-scale feature-aggregation *neck* (PANet/BiFPN), where high-entropy features can be forced through all pyramid levels; the *dual-label assignment* head, where conflicting one-to-one and one-to-many decisions can stall the score comparator; the *task-aligned* top- k selection, where many near-identical scores maximize sorting work; *deformable attention*, where maximizing sampling-point spread induces non-coalesced memory access; and the *one-to-one matching* itself, where near-uniform assignment costs drive worst-case Hungarian iterations ($O(N^3)$). From a real-time-systems perspective, the matching-deadlock surface is especially compelling because it directly threatens worst-case execution-time (WCET) bounds—a well-defined failure mode in safety-critical scheduling. We highlight this as a concrete, under-explored direction: the persistence of latency threats across an architectural change that was *not* motivated by security.

10.2 Cross-cutting Open Challenges

Standardized, system-level benchmarks. Most attacks report per-component timing in incompatible units; few connect to downstream outcomes. SlowTrack’s crash-rate evaluation [60] and CP-FREEZER’s testbed [86] are exceptions. The field needs benchmarks that report *both* timing and safety/serviceability impact on common axes.

Physical realizability. Projector attacks [61, 65] demonstrate multi-second physical delays but require active infrastructure; SlowTrack’s physical transfer remains preliminary [60]. The gap between strong digital results and robust physical deployment is wide.

Inference-time compute as a new surface. Reasoning models and MoE routers [20, 49, 55] created entirely new attack classes within a single year. As architectures continue to make compute input-dependent, each such mechanism is a candidate target, and defenses lag behind by construction.

Timing as an oracle. Black-box timing attacks [66] show that even opaque deployments leak optimization signals, complicating any defense that assumes the attacker’s ignorance.

The efficiency–robustness tradeoff. Every architectural defense—pruning, sparsity, early exit—trades average-case efficiency for worst-case robustness, and the tradeoff is non-monotone [33, 81]. Defenses that bound worst-case cost by construction (budget enforcement, WCET guards) are more promising than those that merely reduce average cost.

Detection fragility. High detection-bypass rates [55] argue that anomaly detection alone is insufficient; enforceable budgets and scheduling must complement it.

End-to-end stack evaluation. A modern autonomy stack with a VLM planner is simultaneously vulnerable at the camera sensor, NMS, MOT, LiDAR/cooperative fusion, and the LLM decoder, yet no published work evaluates all of these jointly. Building a simulation environment that injects attack stimuli at each layer and measures *system-level* outcomes—collision rate, mission completion, service availability—is, in our view, the most practically important open engineering problem in this space, and the natural home for the work-budget abstraction discussed in Section 9.

Defense non-transfer. Because perception-side and serving-side defenses operate on incompatible attack surfaces (Section 9), a heterogeneous stack must currently compose multiple narrowly scoped defenses. Whether a single unifying primitive—work-budget enforcement at every stage boundary—can subsume them without unacceptable accuracy cost is an open question.

Collectively, these challenges indicate that latency attacks are transitioning from isolated algorithmic curiosities to a fundamental systems-security problem. Addressing them will likely require tighter integration between machine learning, systems, hardware, and real-time computing than current research communities typically consider.

11 Conclusion

Latency attacks extend the security of deep learning beyond correctness to timeliness and resource availability. What began as NMS overload on a single detector has matured into a cross-cutting threat spanning physical autonomous-driving perception, sponge attacks on accelerators, slowdown attacks on every form of conditional computation, and a rapidly expanding frontier of output-length, verbose-image, and reasoning denial-of-service attacks on LLMs and VLMs. The unifying lesson is that *any* mechanism which makes inference cost input-dependent—post-processing, sparsity, early exit, routing, autoregressive generation, inference-time reasoning—becomes an attack surface the moment efficiency is advertised. Defenses are advancing, with hardware-aware training, input purification, and especially budget-enforcement-by-construction showing the most promise. Still, they remain fragmented and bound by a persistent efficiency–robustness tradeoff. Progress will depend on standardized benchmarks that report timing alongside downstream impact, on threat models that take physical and serving realism seriously, and on defenses that control worst-case computational paths rather than merely lowering the average. As deep learning continues to trade fixed cost for input-adaptive efficiency, the availability axis will only grow in importance. We expect availability-oriented attacks to become a permanent component of adversarial machine learning research rather than a niche topic. As future architectures increasingly expose dynamic computation, understanding worst-case inference behavior will likely become as important as improving average-case efficiency.

References

- [1] Francisco Aguilera-Martínez and Fernando Berzal Galiano. 2025. LLM Security: Vulnerabilities, Attacks, Defenses, and Countermeasures. arXiv preprint arXiv:2505.01177. doi:10.48550/arXiv.2505.01177
- [2] Muhammad Shakeel Akram, Vincent Meyers, and Mehdi B. Tahoori. 2025. EvoWeight: Sponge Poisoning of FPGA-Based DNN Accelerators in Differentially Private Secure Federated Learning. In *Proceedings of the IEEE International Symposium on Hardware Oriented Security and Trust (HOST)*. IEEE, Piscataway, NJ, USA, 1–11. doi:10.1109/HOST64725.2025.11050058
- [3] Amit Baras, Alon Zolfi, Asaf Shabtai, and Yuval Elovici. 2025. QuantAttack: Exploiting Dynamic Quantization to Attack Vision Transformers. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*. IEEE, Tucson, AZ, USA, 1–11. arXiv:2312.02220 doi:10.48550/arXiv.2312.02220
- [4] Dudi Biton, Aditi Misra, Efrat Levy, Jaidip Kotak, Ron Bitton, Roei Schuster, Nicolas Papernot, Yuval Elovici, and Ben Nassi. 2023. The Adversarial Implications of Variable-Time Inference. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISec@CCS)*. ACM, New York, NY, USA, 103–114. doi:10.1145/3605764.3623912

- [5] Hanene F. Z. Brachemi Meftah, Wassim Hamidouche, Sid Ahmed Fezza, and Olivier Deforges. 2025. Energy-Latency Attacks: A New Adversarial Threat to Deep Learning. arXiv preprint arXiv:2503.04963. doi:10.48550/arXiv.2503.04963
- [6] Nicolas Carion, Francisco Massa, Gabriel Synnaeve, Nicolas Usunier, Alexander Kirillov, and Sergey Zagoruyko. 2020. End-to-End Object Detection with Transformers. In *European Conference on Computer Vision (ECCV)*. Springer, Cham, Switzerland, 213–229. doi:10.1007/978-3-030-58452-8_13
- [7] Erh-Chung Chen, Pin-Yu Chen, I-Hsin Chung, and Che-Rung Lee. 2024. Overload: Latency Attacks on Object Detection for Edge Devices. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, Piscataway, NJ, USA, 24716–24725. doi:10.1109/CVPR52733.2024.02334
- [8] Hanlin Chen, Simin Chen, Wenyu Li, Wei Yang, and Yiheng Feng. 2025. Impact Analysis of Inference Time Attack of Perception Sensors on Autonomous Vehicles. arXiv preprint arXiv:2505.03850. doi:10.48550/arXiv.2505.03850
- [9] Hao Chen, Xin Liu, Jie Jiang, et al. 2026. Token-Budget-Aware Pool Routing for Cost-Efficient LLM Inference. arXiv preprint arXiv:2604.09613. arXiv:2604.09613 doi:10.48550/arXiv.2604.09613
- [10] Simin Chen, Cong Liu, Mirazul Haque, Zihe Song, and Wei Yang. 2022. NMTSlot: Understanding and Testing Efficiency Degradation of Neural Machine Translation Systems. In *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE)*. ACM, New York, NY, USA, 1148–1160. doi:10.1145/3540250.3549102
- [11] Simin Chen, Zihe Song, Mirazul Haque, Cong Liu, and Wei Yang. 2022. NIGSLOWDOWN: Evaluating the Efficiency Robustness of Neural Image Caption Generation Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, Piscataway, NJ, USA, 15365–15374. doi:10.1109/CVPR52688.2022.01493
- [12] Yiming Chen, Simin Chen, Zexin Li, Wei Yang, Cong Liu, Robby T. Tan, and Haizhou Li. 2023. Dynamic Transformers Provide a False Sense of Efficiency. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*. Association for Computational Linguistics, Stroudsburg, PA, USA, 7180–7193. doi:10.48550/arXiv.2305.12228
- [13] Antonio Emanuele Cinà, Ambra Demontis, Battista Biggio, Fabio Roli, and Marcello Pelillo. 2022. Energy-Latency Attacks via Sponge Poisoning. arXiv preprint arXiv:2203.08147. doi:10.48550/arXiv.2203.08147
- [14] Aobotao Dai, Xinyu Ma, Lei Chen, Songze Li, and Lin Wang. 2025. When Data Manipulation Meets Attack Goals: An In-depth Survey of Attacks for VLMs. arXiv preprint arXiv:2502.06390. doi:10.48550/arXiv.2502.06390
- [15] Ali Dehghantanha and Sajad Homayoun. 2026. SoK: The Attack Surface of Agentic AI – Tools and Autonomy. arXiv preprint arXiv:2603.22928. doi:10.48550/arXiv.2603.22928
- [16] Jianshuo Dong, Ziyuan Zhang, Qingjie Zhang, Tianwei Zhang, Hao Wang, Hewu Li, Qi Li, Chao Zhang, Ke Xu, and Han Qiu. 2024. An Ergorgio Prompt Makes Large Language Model Babble on. arXiv preprint arXiv:2412.19394. doi:10.48550/arXiv.2412.19394
- [17] Dongfang Du, Chaofeng Sha, and Xin Peng. 2025. Energy Attack Method for Adaptive Multi-Exit Neural Networks. *Information and Software Technology* 179 (2025), 107653. doi:10.1016/j.infsof.2024.107653
- [18] Anthony Etim and Jakub Szefer. 2024. Time Traveling to Defend Against Adversarial Example Attacks in Image Classification. arXiv preprint arXiv:2410.08338. doi:10.48550/arXiv.2410.08338
- [19] William Fedus, Barret Zoph, and Noam Shazeer. 2022. Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity. arXiv preprint arXiv:2101.03961. arXiv:2101.03961 doi:10.48550/arXiv.2101.03961
- [20] Zekun Fei, Zihao Wang, Weijie Liu, Ruiqi He, Jianing Geng, Zheli Liu, and Xiaofeng Wang. 2026. Misrouter: Exploiting Routing Mechanisms for Input-Only Attacks on Mixture-of-Experts LLMs. arXiv preprint arXiv:2605.04446. doi:10.48550/arXiv.2605.04446
- [21] Xiaoning Feng, Xiaohong Han, Simin Chen, and Wei Yang. 2022. LLMEffChecker: Understanding and Testing Efficiency Degradation of Large Language Models. arXiv preprint arXiv:2210.03696. doi:10.48550/arXiv.2210.03696
- [22] Jiyuan Fu, Kaixun Jiang, Lingyi Hong, Jinglun Li, Haijing Guo, Dingkan Yang, Zhaoyu Chen, and Wenqiang Zhang. 2025. LingoLoop Attack: Trapping MLLMs via Linguistic Context and State Entrapment into Endless Loops. arXiv preprint arXiv:2506.14493. doi:10.48550/arXiv.2506.14493
- [23] Xiaowei Fu and Lei Zhang. 2025. Adversarial Defense in Vision-Language Models: An Overview. In *International Conference on Image Processing, Computer Vision and Machine Learning (ICICML)*. IEEE, Piscataway, NJ, USA, 1–6. doi:10.1109/ICICML67980.2025.11333826
- [24] Matteo Gambella, Fabrizio Pittorino, Giuliano Casale, and Manuel Roveri. 2026. SQUAD: Scalable Quorum Adaptive Decisions via Ensemble of Early Exit Neural Networks. arXiv preprint arXiv:2601.22711. doi:10.48550/arXiv.2601.22711
- [25] Kuofeng Gao, Yang Bai, Jindong Gu, Shu-Tao Xia, Philip Torr, Zhifeng Li, and Wei Liu. 2024. Inducing High Energy-Latency of Large Vision-Language Models with Verbose Images. In *International Conference on Learning Representations (ICLR)*. OpenReview.net, Online, 1–26. doi:10.48550/arXiv.2401.11170
- [26] Kuofeng Gao, Tianyu Pang, Chao Du, Yong Yang, Shu-Tao Xia, and Min Lin. 2024. Denial-of-Service Poisoning Attacks against Large Language Models. arXiv preprint arXiv:2410.10760. doi:10.48550/arXiv.2410.10760

- [27] Xiaoxue Gao, Yiming Chen, Xianghu Yue, Yu Tsao, and Nancy F. Chen. 2025. TTSslow: Slow Down Text-to-Speech with Efficiency Robustness Evaluations. arXiv preprint arXiv:2407.01927. doi:10.48550/arXiv.2407.01927
- [28] Zeyu Gao, Amin Saremi, and Zonghua Gu. 2026. Performance Measurement and Analysis of Certifiable Defenses against Adversarial Patch Attacks. *Journal of Signal Processing Systems* 98, 1 (2026), 22.
- [29] Soumendu Kumar Ghosh, Arnab Raha, and Vijay Raghunathan. 2023. Energy-Efficient Approximate Edge Inference Systems. *ACM Transactions on Embedded Computing Systems* 22, 4 (2023), 77:1–77:50. doi:10.1145/3589766
- [30] Arijit Ghosh Chowdhury, Md. Mofijul Islam, Vaibhav Kumar, F. H. Shezan, Vinija Jain, and Aman Chadha. 2024. Breaking Down the Defenses: A Comparative Survey of Attacks on Large Language Models. arXiv preprint arXiv:2403.04786. doi:10.48550/arXiv.2403.04786
- [31] Shikha Goel, Rajesh Kedia, Rijurekha Sen, and M. Balakrishnan. 2025. EXPRESS: A Framework for Execution Time Prediction of Concurrent CNNs on Xilinx DPU Accelerator. *ACM Transactions on Embedded Computing Systems* 24, 1 (2025), 11:1–11:31. doi:10.1145/3697835
- [32] Mirazul Haque, Anki Chauhan, Cong Liu, and Wei Yang. 2020. ILFO: Adversarial Attack on Adaptive Neural Networks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, Seattle, WA, USA, 14264–14273. doi:10.1109/CVPR42600.2020.01428
- [33] Syed Mhamudul Hasan, Hussein Zangoti, Iraklis Anagnostopoulos, and Abdur R. Shahid. 2025. Sponge Attacks on Sensing AI: Energy-Latency Vulnerabilities and Defense via Model Pruning. In *Proceedings of the IEEE Global Communications Conference (GLOBECOM)*. IEEE, Piscataway, NJ, USA, 1–6. doi:10.1109/GLOBECOM59602.2025.11432163
- [34] Sanghyun Hong, Yiğitcan Kaya, Ionuț-Vlad Modoranu, and Tudor Dumitras. 2021. A Panda? No, It's a Sloth: Slowdown Attacks on Adaptive Multi-Exit Neural Network Inference. In *International Conference on Learning Representations (ICLR)*. OpenReview.net, Online, 1–23. arXiv:2010.02432 doi:10.48550/arXiv.2010.02432
- [35] Md. Iqbal Hossain, Neeresh Kumar Perla, Afia Sajeeda, Siyu Xia, and Ming Shao. 2026. A Survey of Recent Advances in Adversarial Attack and Defense on Vision-Language Models. arXiv preprint. doi:10.1016/j.neunet.2026.108764
- [36] Mengjie Hu, Zirui Li, Haotian Wang, Yu Wang, Yufei Ding, Chun Liu, and Qing Song. 2026. Trajectory-Aware Attack: Explainable Adversarial Attack against Multiple Object Trackers. *IEEE Transactions on Multimedia* 27 (2026), 1–14.
- [37] Benxuan Huang, Lihui Pang, and Anmin Fu. 2024. Sponge Attack against Multi-Exit Networks with Data Poisoning. *IEEE Access* 12 (2024), 33843–33851. doi:10.1109/ACCESS.2024.3370849
- [38] Ruixuan Huang, Qingyue Wang, Hantao Huang, Yudong Gao, Dong Chen, Shuai Wang, and Wei Wang. 2025. RepetitionCurse: Measuring and Understanding Router Imbalance in Mixture-of-Experts LLMs under DoS Stress. arXiv preprint arXiv:2512.23995. doi:10.48550/arXiv.2512.23995
- [39] Yiyang Huang, Zixuan Wang, Zishen Wan, Yapeng Tian, Hao Xu, Yinhe Han, and Yiming Gan. 2025. ANNIE: Be Careful of Your Robots. arXiv preprint arXiv:2509.03383. doi:10.48550/arXiv.2509.03383
- [40] Yunhan Jia, Yantao Lu, Junjie Shen, Qi Alfred Chen, Hao Chen, Zhenyu Zhong, and Tao Wei. 2020. Fooling Detection Alone is Not Enough: Adversarial Attack against Multiple Object Tracking. In *International Conference on Learning Representations (ICLR)*. OpenReview.net, Online, 1–14. arXiv:1905.11026 doi:10.48550/arXiv.1905.11026
- [41] Weiwei Ju, Wei Bao, Liming Ge, and Dong Yuan. 2021. Dynamic Early Exit Scheduling for Deep Neural Network Inference through Contextual Bandits. In *Proceedings of the 30th ACM International Conference on Information and Knowledge Management (CIKM)*. ACM, New York, NY, USA, 823–832. doi:10.1145/3459637.3482335
- [42] Donghua Kang, Dongjoo Kim, Sung-Kil Ko, et al. 2025. Timestep-Compressed Attack on Spiking Neural Networks through Timestep-Level Backpropagation. In *Proceedings of the AAAI Conference on Artificial Intelligence*. AAAI Press, Washington, DC, USA, 1–9. doi:10.48550/arXiv.2508.13812
- [43] Mazaher Kianpour and Shao-Fang Wen. 2019. Timing Attacks on Machine Learning: State of the Art. arXiv preprint, 111–125 pages. doi:10.1007/978-3-030-29516-5_10
- [44] Jiyong Kim, Jaeho Lee, Jiahao Lin, Alish Kanani, Miao Sun, Ümit Y. Ogras, and Jaehyun Park. 2025. eMamba: Efficient Acceleration Framework for Mamba Models in Edge Computing. *ACM Transactions on Embedded Computing Systems* 24, 5s (2025), 110:1–110:22. doi:10.1145/3756875
- [45] Sarada Krithivasan, Sanchari Sen, and Anand Raghunathan. 2020. Sparsity Turns Adversarial: Energy and Latency Attacks on Deep Neural Networks. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* 39, 11 (2020), 4129–4141. doi:10.1109/TCAD.2020.3013077
- [46] Sarada Krithivasan, Sanchari Sen, Nitin Rathi, Kaushik Roy, and Anand Raghunathan. 2022. Efficiency Attacks on Spiking Neural Networks. In *Proceedings of the 59th ACM/IEEE Design Automation Conference (DAC)*. ACM, New York, NY, USA, 373–378. doi:10.1145/3489517.3530443
- [47] Min Li, Xin Fan, Qing Jing, et al. 2025. LDP: Latent Diffusion-based Adversarial Purification towards Transformer-based Visual Encoders. In *Proceedings of the IEEE International Conference on Trust, Security and Privacy in Computing and Communications (TrustCom)*. IEEE, Piscataway, NJ, USA, 1–8. doi:10.1109/Trustcom66490.2025.00065

- [48] Yu Li, Biao Huang, Jinyin Hu, and Cheng Zhuo. 2026. SPLAT: Revisiting Latency Attack on Dynamic Neural Networks. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* 45, 1 (Jan. 2026), 506–518. doi:10.1109/TCAD.2025.3576320
- [49] Yunzhe Li, Jianan Wang, Hongzi Zhu, James Lin, Shan Chang, and Minyi Guo. 2026. ThinkTrap: Denial-of-Service Attacks against Black-box LLM Services via Infinite Thinking. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*. Internet Society, Reston, VA, USA, 1–16. arXiv:2512.07086 doi:10.48550/arXiv.2512.07086
- [50] Jin Liang, Ran Yi, Jie Chen, et al. 2024. Securing Autonomous Vehicles Visual Perception: Adversarial Patch Attack and Defense Schemes with Experimental Validations. *IEEE Transactions on Intelligent Vehicles* 9, 7 (2024), 5234–5247. doi:10.1109/TIV.2024.3403667
- [51] Zi Liang, Ronghua Li, Yanyun Wang, Qingqing Ye, and Haibo Hu. 2026. Can a Single Message Paralyze the AI Infrastructure? The Rise of AbO-DDoS Attacks through Targeted Mobius Injection. arXiv preprint arXiv:2605.11442. doi:10.48550/arXiv.2605.11442
- [52] Zhiyu Liao, Kang Chen, Y. Lin, Kangkang Li, Yunxuan Liu, Hefeng Chen, Xingwang Huang, and Yuanhui Yu. 2025. Attack and Defense Techniques in Large Language Models: A Survey and New Perspectives. arXiv preprint arXiv:2505.00976. doi:10.48550/arXiv.2505.00976
- [53] Han Liu, Yuhao Wu, Zhiyuan Yu, Yevgeniy Vorobeychik, and Ning Zhang. 2023. SlowLiDAR: Increasing the Latency of LiDAR-Based Detection Using Adversarial Examples. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, Piscataway, NJ, USA, 5146–5155. doi:10.1109/CVPR52729.2023.00498
- [54] Ming Liu, Siyuan Liang, Koushik Howlader, Liwen Wang, Dacheng Tao, and Wensheng Zhang. 2025. Natural Reflection Backdoor Attack on Vision Language Model for Autonomous Driving. arXiv preprint arXiv:2505.06413. doi:10.48550/arXiv.2505.06413
- [55] Xiaogeng Liu, Xinyan Wang, Yechao Zhang, Sanjay Kariyappa, Chong Xiang, Muhao Chen, G. Edward Suh, and Chaowei Xiao. 2026. ReasoningBomb: A Stealthy Denial-of-Service Attack by Inducing Pathologically Long Reasoning in Large Reasoning Models. In *Proceedings of the ACM Conference on Computer and Communications Security (CCS)*. ACM, New York, NY, USA, 1–15. doi:10.48550/arXiv.2602.00154
- [56] Tobias Lorenz, Marta Kwiatkowska, and Mario Fritz. 2021. Certifiers Make Neural Networks Vulnerable to Availability Attacks. arXiv preprint arXiv:2108.11299. doi:10.48550/arXiv.2108.11299
- [57] Xiangzhong Luo, Di Liu, Hao Kong, Shuo Huai, Hui Chen, Guochu Xiong, and Weichen Liu. 2024. Efficient Deep Learning Infrastructures for Embedded Computing Systems: A Comprehensive Survey and Future Envision. *ACM Transactions on Embedded Computing Systems* 24, 1 (2024), 1–100. doi:10.1145/3701728
- [58] Zhi Luo, Zenghui Yuan, Wenqi Wei, Daizong Liu, and Pan Zhou. 2025. An Image Is Worth Ten Thousand Words: Verbose-Text Induction Attacks on VLMs. arXiv preprint arXiv:2511.16163. doi:10.48550/arXiv.2511.16163
- [59] Chenxiang Ma, Ningfei Wang, Qi Alfred Chen, and Chao Shen. 2023. WIP: Towards the Practicality of the Adversarial Attack on Object Tracking in Autonomous Driving. In *Symposium on Vehicle Security and Privacy (VehicleSec)*. Internet Society, Reston, VA, USA, 1–8. doi:10.14722/vehiclesec.2023.23063
- [60] Chen Ma, Ningfei Wang, Qi Alfred Chen, and Chao Shen. 2024. SlowTrack: Increasing the Latency of Camera-Based Perception in Autonomous Driving Using Adversarial Examples. In *Proceedings of the AAAI Conference on Artificial Intelligence*. AAAI Press, Washington, DC, USA, 4062–4070. doi:10.1609/aaai.v38i5.28200
- [61] Chen Ma, Ningfei Wang, Zhengyu Zhao, Qi Alfred Chen, and Chao Shen. 2024. SlowPerception: Physical-World Latency Attack against Visual Perception in Autonomous Driving. arXiv preprint arXiv:2406.05800. doi:10.48550/arXiv.2406.05800
- [62] Vicent Sanz Marco, Ben Taylor, Zheng Wang, and Yehia Elkhatib. 2020. Optimizing Deep Learning Inference on Embedded Systems Through Adaptive Model Selection. *ACM Transactions on Embedded Computing Systems* 19, 1 (2020), 2:1–2:28. doi:10.1145/3371154
- [63] Jiyang Mu. 2024. A Real-Time Defense against Object Vanishing Adversarial Patch Attacks for Object Detection in Autonomous Vehicles. arXiv preprint arXiv:2412.06215. doi:10.48550/arXiv.2412.06215
- [64] Andreas Müller and Erwin Quiring. 2024. The Impact of Uniform Inputs on Activation Sparsity and Energy-Latency Attacks in Computer Vision. In *Proceedings of the IEEE Security and Privacy Workshops (SPW)*. IEEE, Piscataway, NJ, USA, 104–111. doi:10.1109/SPW63631.2024.00016
- [65] Raymond Muller, Ruoyu Song, Chenyi Wang, Yanmao Zhao, Hyungjoon Kim, Qi Alfred Chen, and Z. Berkay Celik. 2025. Investigating Physical Latency Attacks against Camera-Based Perception. In *Proceedings of the IEEE Symposium on Security and Privacy (S&P)*. IEEE, Piscataway, NJ, USA, 3782–3800. doi:10.1109/SP61157.2025.00236
- [66] Tsunato Nakai, Daisuke Suzuki, Fumio Omatsu, and Takeshi Fujino. 2021. Adversarial Black-Box Attacks with Timing Side-Channel Leakage. *IEICE Transactions on Fundamentals of Electronics, Communications and Computer Sciences* 104, 1 (2021), 143–151. doi:10.1587/transfun.2020CIP0022
- [67] K. L. Navaneet, Soroush Abbasi Koohpayegani, Essam Sleiman, and Hamed Pirsiavash. 2024. SlowFormer: Adversarial Attack on Compute and Energy Consumption of Efficient Vision Transformers. In *Proceedings of the IEEE/CVF*

- Conference on Computer Vision and Pattern Recognition (CVPR). IEEE, Seattle, WA, USA, 24786–24797. arXiv:2310.02544 doi:10.48550/arXiv.2310.02544
- [68] OWASP Foundation. 2025. *LLM10:2025 Unbounded Consumption*. Technical Report. OWASP Top 10 for Large Language Model Applications. Version 2025. <https://owasp.org/www-project-top-10-for-large-language-model-applications/>
- [69] Souvik Paul and Nicolas Kourtellis. 2023. Poster: Sponge ML Model Attacks of Mobile Apps. In *Proceedings of the 24th International Workshop on Mobile Computing Systems and Applications (HotMobile)*. ACM, New York, NY, USA, 139–139. doi:10.1145/3572864.3581586
- [70] Yongming Rao, Wenliang Zhao, Benlin Liu, Jiwen Lu, Jie Zhou, and Cho-Jui Hsieh. 2021. DynamicViT: Efficient Vision Transformers with Dynamic Token Sparsification. In *Advances in Neural Information Processing Systems (NeurIPS)*. Curran Associates, Red Hook, NY, USA, 13937–13949. arXiv:2106.02034 doi:10.48550/arXiv.2106.02034
- [71] Ravishka Rathnasuriya and Wei Yang. 2025. Exploiting Efficiency Vulnerabilities in Dynamic Deep Learning Systems. arXiv preprint arXiv:2506.17621. doi:10.48550/arXiv.2506.17621
- [72] Yu Ren, Qiang Li, Zhuo Xiong, et al. 2025. DCT-based Fast Localization and Elimination of Adversarial Patches. In *International Conference on Real-time Computing and Robotics (RCAR)*. IEEE, Piscataway, NJ, USA, 1–7. doi:10.1109/RCAR65431.2025.11139457
- [73] Debopam Sanyal, Jui-Tse Hung, Manavi Agrawal, Prahlad Jasti, Shahab Nikkhoo, Somesh Jha, Tianhao Wang, Sibin Mohan, and Alexey Tumanov. 2023. Pareto-Secure Machine Learning (PSML): Fingerprinting and Securing Inference Serving Systems. arXiv preprint arXiv:2307.01292. doi:10.48550/arXiv.2307.01292
- [74] Coen Schoof, Stefanos Koffas, Mauro Conti, and Stjepan Picek. 2024. Beyond PhantomSponges: Enhancing Sponge Attack on Object Detection Models. In *Proceedings of the 2024 ACM Workshop on Wireless Security and Machine Learning (WiseML)*. ACM, New York, NY, USA, 14–19. doi:10.1145/3649403.3656485
- [75] Avishag Shapira, Alon Zolfi, Luca Demetrio, Battista Biggio, and Asaf Shabtai. 2023. Phantom Sponges: Exploiting Non-Maximum Suppression to Attack Deep Object Detectors. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*. IEEE, Piscataway, NJ, USA, 4571–4580. doi:10.1109/WACV56688.2023.00455
- [76] Ilia Shumailov, Yiren Zhao, Daniel Bates, Nicolas Papernot, Robert Mullins, and Ross Anderson. 2021. Sponge Examples: Energy-Latency Attacks on Neural Networks. In *Proceedings of the IEEE European Symposium on Security and Privacy (EuroS&P)*. IEEE, Piscataway, NJ, USA, 212–231. doi:10.1109/EuroSP51992.2021.00024
- [77] Manojna Sistla, Yu Wen, Aamir Bader Shah, Chenpei Huang, Lening Wang, Xuqing Wu, et al. 2025. Bit-Flip Induced Latency Attacks in Object Detection. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*. IEEE, Piscataway, NJ, USA, 6651–6660. doi:10.1109/WACV61041.2025.00653
- [78] Tarun Aditya Srikrishnan, Devang Shah, and Steven K. Reinhardt. 2025. BlindSight: Harnessing Sparsity for Efficient VLMs. arXiv preprint arXiv:2507.09071. doi:10.48550/arXiv.2507.09071
- [79] Jonathan Steinberg and Oren Gal. 2026. Semantic Denial of Service in LLM-Controlled Robots. arXiv preprint arXiv:2604.24790. doi:10.48550/arXiv.2604.24790
- [80] Duoxun Tang, Dasen Dai, Jiyao Wang, Xiao Yang, Jianyu Wang, and Siqi Cai. 2026. VidDoS: Universal Denial-of-Service Attack on Video-based Large Language Models. arXiv preprint arXiv:2603.01454. doi:10.48550/arXiv.2603.01454
- [81] Jona te Lintelo, Stefanos Koffas, and Stjepan Picek. 2024. The SkipSponge Attack: Sponge Weight Poisoning of Deep Neural Networks. arXiv preprint arXiv:2402.06357. doi:10.48550/arXiv.2402.06357
- [82] Surat Teerapittayanon, Bradley McDanel, and H. T. Kung. 2016. BranchyNet: Fast Inference via Early Exiting from Deep Neural Networks. In *Proceedings of the 23rd International Conference on Pattern Recognition (ICPR)*. IEEE, Piscataway, NJ, USA, 2464–2469. doi:10.1109/ICPR.2016.7900006
- [83] Kamala Varma, Arda Numanoglu, Yigitcan Kaya, and Tudor Dumitras. 2024. Understanding, Uncovering, and Mitigating the Causes of Inference Slowdown for Language Models. In *Proceedings of the IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*. IEEE, Piscataway, NJ, USA, 723–740.
- [84] Apostol Vassilev, Alina Oprea, Alie Fordyce, and Hyrum Anderson. 2023. *Adversarial Machine Learning: A Taxonomy and Terminology of Attacks and Mitigations*. Technical Report NIST AI 100-2e2023. National Institute of Standards and Technology. doi:10.6028/NIST.AI.100-2e2023
- [85] Ao Wang, Hui Chen, Lihao Liu, Kai Chen, Zijia Lin, Jungong Han, and Guiguang Ding. 2024. YOLOv10: Real-Time End-to-End Object Detection. In *Advances in Neural Information Processing Systems (NeurIPS)*. Curran Associates, Red Hook, NY, USA, 107984–108011. arXiv:2405.14458 doi:10.48550/arXiv.2405.14458
- [86] Chenyi Wang, Ruoyu Song, Raymond Muller, Jean-Philippe Monteui, Z. Berkay Celik, Jonathan Petit, and Qi Alfred Chen. 2026. CP-FREEZER: Latency Attacks against Vehicular Cooperative Perception. In *Proceedings of the AAAI Conference on Artificial Intelligence*. AAAI Press, Washington, DC, USA, 1–9. doi:10.1609/aaai.v40i2.37082
- [87] Derui Wang, Chaoran Li, Sheng Wen, Surya Nepal, and Yang Xiang. 2019. Daedalus: Breaking Non-Maximum Suppression in Object Detection via Adversarial Examples. arXiv preprint arXiv:1902.02067. arXiv:1902.02067 doi:10.48550/arXiv.1902.02067

- [88] Tianyi Wang, Huawei Fan, Yuanchao Shu, Peng Cheng, and Cong Wang. 2026. Rethinking Latency Denial-of-Service: Attacking the LLM Serving Framework, Not the Model. arXiv preprint arXiv:2602.07878. doi:10.48550/arXiv.2602.07878
- [89] Tianyi Wang, Zichen Wang, Cong Wang, Yuanchao Shu, Ruilong Deng, Peng Cheng, and Jiming Chen. 2025. Can't Slow Me Down: Learning Robust and Hardware-Adaptive Object Detectors against Latency Attacks for Edge Devices. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. IEEE, Piscataway, NJ, USA, 19090–19100. doi:10.1109/CVPR52734.2025.01791
- [90] Xin Wang, Jie Li, Zejia Weng, Yixu Wang, Yifeng Gao, Tianyu Pang, Chao Du, Yan Teng, Yingchun Wang, Zuxuan Wu, Xingjun Ma, and Yu-Gang Jiang. 2025. FreezeVLA: Action-Freezing Attacks against Vision-Language-Action Models. arXiv preprint arXiv:2509.19870. doi:10.48550/arXiv.2509.19870
- [91] Xin Wang, Anand Suresh, Alan Zhang, et al. 2026. Conformal Thinking: Risk Control for Reasoning on a Compute Budget. arXiv preprint arXiv:2602.03814. doi:10.48550/arXiv.2602.03814
- [92] Xiasi Wang, Tianliang Yao, Simin Chen, Runqi Wang, Lei Ye, Kuofeng Gao, Yi Huang, and Yuan Yao. 2025. VLMInfer-Slow: Evaluating the Efficiency Robustness of Large Vision-Language Models as a Service. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*. Association for Computational Linguistics, Stroudsburg, PA, USA, 16009–16024. doi:10.18653/v1/2025.acl-long.781
- [93] Yan Wang, Xingbin Wang, Zechao Lin, Yulan Su, Sisi Zhang, Rui Hou, and Dan Meng. 2024. Garrison: A High-Performance GPU-Accelerated Inference System for Adversarial Ensemble Defense. In *Proceedings of the 61st ACM/IEEE Design Automation Conference (DAC)*. ACM, New York, NY, USA, 1–6. doi:10.1145/3649329.3654810
- [94] Zijian Wang, Shuo Huang, Yujin Huang, and Helei Cui. 2023. Energy-Latency Attacks to On-Device Neural Networks via Sponge Poisoning. In *Proceedings of the 2023 ACM Workshop on Secure and Trustworthy Deep Learning Systems*. ACM, New York, NY, USA, 1–11. doi:10.1145/3591197.3591307
- [95] Lichen Xia, Xing Gao, and Weisong Shi. 2026. Investigating Stealthy Latency Attacks Against Camera-Based Object Detection. In *Proceedings of the 4th USENIX Symposium on Vehicle Security and Privacy (VehicleSec)*. USENIX Association, Baltimore, MD, USA, 1–7.
- [96] Yong Xiao, Jin Ma, Ping Yi, and Xiuzhen Chen. 2024. Sponge Backdoor Attack: Increasing the Latency of Object Detection Exploiting Non-Maximum Suppression. In *Proceedings of the International Joint Conference on Neural Networks (IJCNN)*. IEEE, Piscataway, NJ, USA, 1–8. doi:10.1109/IJCNN60899.2024.10650435
- [97] Wenpeng Xing, Minghao Li, Mohan Li, and Meng Han. 2025. Towards Robust and Secure Embodied AI: A Survey on Vulnerabilities and Attacks. arXiv preprint arXiv:2502.13175. doi:10.48550/arXiv.2502.13175
- [98] Haochuan Xu, Yun Sing Koh, Shuhuai Huang, Zirun Zhou, Di Wang, Jun Sakuma, and Jingfeng Zhang. 2025. Model-Agnostic Adversarial Attack and Defense for Vision-Language-Action Models. arXiv preprint arXiv:2510.13237. doi:10.48550/arXiv.2510.13237
- [99] Wenrui Xu and Keshab K. Parhi. 2025. A Survey of Attacks on Large Language Models. arXiv preprint arXiv:2505.12567. doi:10.48550/arXiv.2505.12567
- [100] Zhiyuan Xu, Joseph Gardiner, Sana Belguith, and Lichao Wu. 2026. RouteHijack: Routing-Aware Attack on Mixture-of-Experts LLMs. arXiv preprint arXiv:2605.02946. doi:10.48550/arXiv.2605.02946
- [101] Shengyao Zhang, Xudong Pan, Mi Zhang, and Min Yang. 2023. SlowBERT: Slow-down Attacks on Input-adaptive Multi-exit BERT. In *Findings of the Association for Computational Linguistics: ACL 2023*. Association for Computational Linguistics, Stroudsburg, PA, USA, 9992–10007. doi:10.18653/v1/2023.findings-acl.634
- [102] Shengyao Zhang, Mi Zhang, Xudong Pan, and Min Yang. 2023. No-Skim: Towards Efficiency Robustness Evaluation on Skimming-based Language Models. arXiv preprint arXiv:2312.09494. doi:10.48550/arXiv.2312.09494
- [103] Yan Zhang, Zihao Liu, Yi Zhu, and Chenglin Miao. 2025. Towards Real-Time Defense against Object-Based LiDAR Attacks in Autonomous Driving. In *Proceedings of the ACM Conference on Computer and Communications Security (CCS)*. ACM, New York, NY, USA, 1–15. doi:10.1145/3719027.3765227
- [104] Yu Zhang, Xiao Wang, Hao Gao, et al. 2025. PD3F: A Pluggable and Dynamic DoS-Defense Framework against Resource Consumption Attacks Targeting Large Language Models. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, Stroudsburg, PA, USA, 1–16. doi:10.48550/arXiv.2505.18680
- [105] Yuanhe Zhang, Zhenhong Zhou, Wei Zhang, Xinyue Wang, Xiaojun Jia, Yang Liu, and Sen Su. 2025. Crabs: Consuming Resource via Auto-generation for LLM-DoS Attack under Black-box Settings. In *Findings of the Association for Computational Linguistics: ACL 2025*. Association for Computational Linguistics, Stroudsburg, PA, USA, 1–18. doi:10.48550/arXiv.2412.13879
- [106] Zhihui Zhao, Xiaorong Dong, Yaowen Zheng, Xiaohui Chen, Yimo Ren, Hangbei Cheng, Yongle Chen, and Limin Sun. 2026. Breaking Cross-Modal Alignment in Embodied Intelligence: A Multimodal Adversarial Attack Framework for Vision-Language-Action Models. In *Proceedings of the ACM Web Conference 2026 (WWW)*. ACM, Sydney, Australia, 1–12. doi:10.1145/3774904.3792315

[107] Kaiyu Zhou, Yongsen Zheng, Yicheng He, Meng Xue, Xueluan Gong, Yuji Wang, Xuanye Zhang, and Kwok-Yan Lam. 2026. Beyond Max Tokens: Stealthy Resource Amplification via Tool Calling Chains in LLM Agents. arXiv preprint arXiv:2601.10955. doi:10.48550/arXiv.2601.10955

[108] Yuguang Zhou, Xunguang Wang, Pingchuan Ma, Zhantong Xue, Zhaoyu Wang, and Shuai Wang. 2026. From Shield to Target: Denial-of-Service Attacks on LLM-Based Agent Guardrails. arXiv preprint arXiv:2606.14517. doi:10.48550/arXiv.2606.14517

A Catalog of Latency, Energy, and Timing Attacks and Defenses

This appendix provides a structured, citable catalog of the works surveyed, organized by attack stage (inference vs. training) and by defenses. Each entry records the target architecture, application domain, attacker setting, and links to the paper and released code where available. A companion website is available here: <https://github.com/guzonghua/awesome-latency-attacks/>. Tables 6–8 mirror that resource.

Table 6. Inference-stage latency, energy, and timing attacks on deep learning.  links to the paper;  links to released code (X = none located).

Attack	Venue (Year)	Target	Domain	Setting	Paper	Code
Daedalus	arXiv (2019)	Object detection (NMS)	CV / AD	White box (physical)		
Sparsity Attacks	IEEE TCAD (2020)	CNNs	CV	White box		X
ILFO	CVPR (2020)	Depth-adaptive networks (AdNNs)	CV	White box		X
Tracker Hijacking	ICLR (2020)	Multi-object tracking	AD	White box		X
Sponge Examples	EuroS&P (2021)	Transformers + CNNs	CV & NLP	Black & white box		
DeepSloth	ICLR (2021)	Multi-exit CNNs	CV	White box		
Timing Side-Channel AE	IEICE Trans. (2021)	DNN classifiers (timing oracle)	CV	Black box (timing)		X
SpikeAttack	ACM/IEEE DAC (2022)	Spiking neural networks	CV	White box		X
NMTSloth	ESEC/FSE (2022)	Decoder-based NMT	NLP	White box		
LLMEffiChecker	ACM TOSEM (2022)	LLMs	NLP	Black & white box		
NICGSlowDown	CVPR (2022)	Decoder-based image captioning	CV	White box		
SAME	ACL (2023)	Multi-exit transformers	NLP	White box		
SlowBERT	ACL Findings (2023)	Multi-exit BERT	NLP	White box		X
No-Skim	arXiv (2023)	Skimming language models	NLP	Black & white box		X
Phantom Sponges	WACV (2023)	Object detection (NMS)	CV / AD	White box		
SlowLiDAR	CVPR (2023)	3D LiDAR detection	CV / AD	White box		
WIP Tracker Attack	VehicleSec (2023)	Multi-object tracking	AD	White box (patch)		X
Variable-Time Inference	AISeC@CCS (2023)	Object detection (NMS timing)	CV	Black box (timing)		X
Overload	CVPR (2024)	Object detection (edge)	CV / AD	White box		X
SlowFormer	CVPR (2024)	Efficient vision transformers	CV	White box (patch)		
Beyond Phantom-Sponges	ACM WiseML (2024)	Object detection (NMS)	CV / AD	White box		X

continued on next page

Table 6. Inference-stage latency, energy, and timing attacks on deep learning.  links to the paper;  links to released code (X = none located). (continued)

Attack	Venue (Year)		Target	Domain	Setting	Paper	Code
SlowTrack	AAAI (2024)		Camera perception (detect+track)	AD	White box		
SlowPerception	arXiv (2024)		Camera perception (NMS+MOT)	AD	White box (physical, projector)		X
TPA (Time-aware)	Pattern Recognition	Recognition Letters	Detection + segmentation	AD	White box		X
Steal Now Attack Later	arXiv (2024)		Object detection	CV	Black box		X
Energy Attack (multi-exit)	Info. & Software Tech. (2024)		Adaptive multi-exit networks	CV	Grey box		X
Slowdown Causes (SaTML)	IEEE SaTML (2024)		Language models	NLP	White box		X
Engorgio	arXiv (2024)		LLMs (output inflation)	NLP	White box + transfer		X
Verbose Images	ICLR (2024)		Large VLMs	CV+NLP	White box		
Uniform Inputs	IEEE SPW (2024)		CNNs (sparsity)	CV	Black box		
DetStorm	IEEE S&P (2025)		Camera perception	AD	White box (physical)		X
Inference-Time Impact Analysis	arXiv (2025)		Full perception (sim.)	AD	Simulation		X
DDLSE Efficiency At- tacks	arXiv (2025)		Early-exit / token-pruning / MoE	CV & NLP	White & black box		X
TTSlow	IEEE TASLP (2025)		Auto-regressive TTS	Speech	White box		X
Crabs	ACL Findings		LLMs (DoS)	NLP	Black box		X
VLMInferSlow	ACL (2025)		VLMs-as-a-service	CV+NLP	Black box		X
Verbose-Text Induc- tion	arXiv (2025)		VLMs	CV+NLP	White box		X
LingoLoop	arXiv (2025)		Multimodal LLMs	CV+NLP	White box		X
VidDoS	arXiv (2026)		Video-LLMs (streaming AD)	CV+NLP / AD	White box (universal patch)		X
FreezeVLA	arXiv (2025)		Vision-language-action models	Robotics	White box (image)		X
Semantic-DoS	arXiv (2026)		LLM-controlled robots	Robotics	Audio injection		X
EDPA	arXiv (2025)		Vision-language-action models	Robotics	Black box (patch)		
MAVLA	ACM WWW (2026)		Vision-language-action models	Robotics	White box		X
ANNIE	arXiv (2025)		Embodied AI / VLA robots	Robotics	White box		
Bit-Flip NMS Attack	WACV (2025)		Object detection (parameters)	CV / AD	Hardware (Rowhammer)		X
QuantAttack	WACV (2025)		Dynamically quantized ViTs	CV	White box		
Timestep- Compressed Attack	AAAI (2025)		Spiking neural networks	CV	White box		X

continued on next page

Table 6. Inference-stage latency, energy, and timing attacks on deep learning.  links to the paper;  links to released code (X = none located). (continued)

Attack	Venue (Year)	Target	Domain	Setting	Paper	Code
CP-FREEZER	AAAI (2026)	Cooperative perception (V2V)	AD	White box (testbed)		X
Trajectory-Aware Attack	IEEE TMM (2026)	Multi-object tracking	AD	White box		X
SPLAT	IEEE TCAD (2026)	Multi-exit dynamic networks	CV	Black box		X
RouteHijack	arXiv (2026)	Mixture-of-experts LLMs	NLP	White box		X
Misrouter	arXiv (2026)	Mixture-of-experts LLMs	NLP	Black box (input-only)		X
ReasoningBomb	ACM CCS (2026)	Large reasoning models	NLP	Black box		X
ThinkTrap	NDSS (2026)	Reasoning LLM APIs	NLP	Black box		X
RepetitionCurse	arXiv (2025)	MoE LLMs (expert parallelism)	NLP	Black box		X
Fill and Squeeze	arXiv (2026)	LLM serving scheduler	NLP	Black box (system)		X
Beyond Max Tokens	arXiv (2026)	LLM agent tool chains (MCP)	NLP / Agentic	Black box		X
From Shield to Target	arXiv (2026)	LLM agent guardrails	NLP / Agentic	Black box (transfer)		X
Mobius Injection (AbO-DDoS)	arXiv (2026)	LLM agent infrastructure	Agentic	Injection		X

Table 7. Training-stage (poisoning / backdoor / weight) latency and energy attacks.

Attack	Venue (Year)	Target	Domain	Setting	Paper	Code
Sponge Poisoning	arXiv / AsiaCCS (2022)	CNNs	CV	Partial control		
On-Device Sponge Poisoning	ACM SecTL (2023)	On-device DNNs	CV	Partial control		X
Mobile-App Sponge	ACM HotMobile (2023)	Mobile ML models	CV	Partial control		X
SkipSponge	arXiv (2024)	CNNs, GANs (weights)	CV	Full control		X
Huang et al. (multi-exit)	IEEE Access (2024)	Multi-exit CNNs	CV	Full control		X
Sponge Backdoor (OD)	IJCNN (2024)	Object detection (NMS)	CV / AD	Backdoor		X
DoS Poisoning (LLM)	arXiv (2024)	LLMs (no-EOS)	NLP	Backdoor / poisoning		X
Sensing-AI Sponge	IEEE GLOBECOM (2025)	Sensing DNNs (IoT)	Sensing	Partial control		X
EvoWeight (FPGA)	IEEE HOST (2025)	FPGA DNN accelerators	CV	Full control		X
Reflection Backdoor (VLM-AD)	arXiv (2025)	Driving VLM planner	AD	Backdoor (physical trigger)		X

Table 8. Defenses against latency, energy, and timing attacks, by control mechanism.

Defense	Venue (Year)	Target	Mechanism	Domain	Paper	Code
DEE Scheduling	ACM CIKM (2021)	Early-exit networks	Runtime scheduling	CV		✗
Certifier Caveat	arXiv (2021)	Certified classifiers	Defense pitfall	CV		✗
Constant-Time NMS	AISeC@CCS (2023)	Object detection (NMS)	Bounded execution	CV		✗
PSML	arXiv (2023)	Inference serving systems	System / serving control	ML serving		✗
Adaptive Resizing	ACM WiseML (2024)	Object detection (NMS)	Input transformation	AD		✗
ADAV Patch Defense	arXiv (2024)	Object detection	Input transformation	AD		✗
Securing AV Perception	IEEE TIV (2024)	AV visual perception	Input transformation	AD		✗
Calibration Defense	IEEE SaTML (2024)	Multi-exit models	Robust training	NLP		✗
Sparsity Monitor	IEEE SPW (2024)	CNNs	Runtime monitoring	CV		
Garrison	ACM/IEEE DAC (2024)	Ensemble inference (GPU)	System / serving control	CV		✗
Time-Traveling Defense	arXiv (2024)	Traffic-sign classifiers	Temporal redundancy	AD		✗
Can't Slow Me Down	CVPR (2025)	Edge object detectors	Robust/adaptive training	AD		✗
DCT Patch Elimination	IEEE RCAR (2025)	Object detection	Input transformation	AD		✗
Real-Time LiDAR Defense	ACM CCS (2025)	LiDAR detection	Runtime monitoring	AD		✗
LDP Purification	IEEE TrustCom (2025)	VLM visual encoders	Input purification	CV+NLP		✗
Pruning Defense	IEEE GLOBECOM (2025)	Sensing DNNs	Architectural sparsity	Sensing		✗
BlindSight	arXiv (2025)	VLMs	Architectural sparsity	CV+NLP		✗
PD3F	EMNLP (2025)	LLM serving	Budget enforcement	NLP		✗
SQUAD	arXiv (2026)	Early-exit ensembles	Runtime scheduling	CV		✗
Token-Budget Routing	arXiv (2026)	LLM serving	Budget enforcement	NLP	✗	✗
Conformal Thinking	arXiv (2026)	Reasoning models	Budget enforcement	NLP		✗